Adaptive Resolution: stage the learning investment

22 October 2026 · Investment recommendation · Synthetic exercise · USD

Recommendation and consumer problem

Approve a limited investment in two-way resolution conversations with focused authenticated tasks. Cap spend at $50,000, with a week-one prerequisite and delivery checkpoint before committing the balance. A four-week launch is conditional on inbound SMS, security and US content approval. Otherwise run a smaller diagnostic or defer production.

Engaged consumers without an arrangement may need suitable timing, account clarification, help or time to decide. Interest is not ability to pay or consent. A conversation can capture the stated need; a focused task can offer an approved action.

Evidence and strongest alternative

The exclusive click, authenticated-visit and incomplete-plan cohorts contain **18,000 consumers** per intake. They already produced 4,600 payers and **$820,000 settled cash** in 30 days under the existing journey. These are fixed cohorts, not a funnel or causal evidence. US scale and SMS reach are unknown.

Eight incomplete-plan interviews reported dates in three, interruption in two, account questions in two and budget concerns in one. Of 50 selected support conversations, 12 concerned dates, nine account questions and seven resumption. These suggest needs, not prevalence. The biggest falsifying assumption is that relevance solves enough barriers; affordability or creditor terms may dominate.

Portal-only guidance is the strongest alternative, using existing authentication with less complexity. Outbound orchestration can aid re-engagement but remains one-way. Conversation exposes explicit needs, yet inbound routing, secure task links and persistence need building. Cost the alternatives in week one.

Commercial sensitivity, not a forecast

Full reach and +3 points imply $94,800 extra cash per intake and about $148,000 annual contribution before new costs. The original 60%-reach +1 to +5-point cases yield $14,026 to $132,336 annually before build. Three points is a reference, not an estimate.

The separate conversational bridge assumes 60% SMS reach, 30% replies, four extra SMS and four AI turns per responder, hypothetical $0.01 per turn, two extra support conversations per 100 responders and zero savings. The supplied $0.012 Composer attempt is not a conversation price. Retain 15% fees, 2% cash costs and $6 support.

Lift among responders	Extra cash per intake	Annual contribution	Simple payback
+1 point	$5,688	-$2,791	None
+3 points	$17,064	$14,956	40.1 intakes
+5 points	$28,440	$32,702	18.3 intakes
+10 points, speculative	$56,880	$77,069	7.8 intakes

This range does not justify rollout at the assumed cost. No lift loses $11,664 annually before build. Reach, responder mix, marginal cash, oversight and infrastructure remain unknown. Twelve new intakes and historical $170/$180/$190 cash per payer are assumptions. Payback excludes build and ramp. Model.

Smaller pilot and autonomy boundaries

The local demo includes deterministic intent classification, remembered preferences, task links and ephemeral drafts. Messages, identity checks, source services, support and payments are simulated. These are proposed capabilities. A plan-start event contains no saved form data; phone-link possession does not verify identity.

The first production pilot covers a small approved intent set, template replies, existing approved options and human help. Reuse existing authentication where possible. Add minimal resume only if feasible. Exclude unapproved generated US content, new creditor terms and autonomous agreements/payments. Consumers explicitly confirm; plans earn no fee until cash settles.

Requested service replies sit outside the proactive budget, not outside all permission constraints. Human compliance must approve their classification and handling. Before proactive handoff, re-read scope, permissions, payment, plan, holds, verified local time and shared cross-channel limits. Unknown state blocks outreach. Apply one attempt per rolling 24 hours and three per seven days, inclusive boundaries, weekdays 09:00 inclusive to 18:00 exclusive. Preserve disclosures. The classifier routes hardship/disputes; humans or account services apply authoritative holds.

Four-week plan and measurement

Week 1: confirm inbound SMS, identity/security, retention, support, approved intents/content and delivery cost. Name compliance and stop owners. Proceed only if prerequisites fit; otherwise use the constrained alternative. Week 2: limited inbound handling and policy-safe routing. Week 3: focused authenticated tasks, instrumentation and minimal resume if feasible. Week 4: integration testing, human compliance approval and a capped launch if prerequisites pass. Four weeks is not mature efficacy evidence.

Randomise consumers 1:1 at t0, stratified by cohort, against the existing journey. Retain non-responders and later suppressed accounts in their assigned arm. Primary outcome is treatment-control 30-day net settled cash per randomised consumer, valued at 13% less measured incremental costs. Responders select themselves; their outcomes cannot establish causal lift. Report uncertainty, exposure, replies, support, complaints and opt-outs. Paid conversion and plans are secondary; follow 60/90-day cash and adherence to distinguish acceleration from durable recovery.

Cash variance is unavailable; size the primary outcome before claiming power. The earlier roughly 18,900-person conversion calculation assumes a 1.8-point assigned effect and does not power this responder funnel. Complete 30-day windows plus next-day reporting before efficacy review.

Expand, change or stop

Expand only with safe, mature randomised evidence whose conservative contribution estimate supports payback within 12 months at verified scale and costs. Validate durability and compare the simpler alternative. Change if reach is weak, support is overloaded or terms dominate. If uncertainty spans the hurdle, permit one extension with a fixed sample, budget and decision date.

Pause for prohibited contact, disclosure failure, unauthorised terms/payment or serious harm. Investigate complaint increases above 0.10 points or opt-out increases above 0.50 points versus control; these proposed triggers do not establish safety below them. Stop for recurring failures, uncontrolled harm or a precise upper contribution estimate below the investment hurdle. AI assisted code, analysis, drafting and checks; the candidate owns the recommendation.