Adaptive Resolution: invest in cohort progression

22 October 2026 · Synthetic exercise · USD

Recommendation and consumer problem

Fund a bounded test for the original 18,000 engaged consumers, capped at $50,000. In week one, confirm feasibility and cost a narrower alternative. Four weeks is a conditional launch horizon. Conversation improves progression only if the experiment establishes it.

Consumers who clicked, visited or started a plan may need suitable dates, account clarification, help or time to decide. A contextual conversation and focused task can help them take the next useful step. It cannot create affordability, change approved terms or treat engagement as permission or commitment.

Cohorts and evidence

The source cohorts are 10,000 click-only consumers, 6,000 portal visitors without a plan start and 2,000 incomplete plans. Only 8,000 had reached the portal. Cohorts are exclusive at t0, not funnel transitions. All were eligible for an allowed contact route at entry.

They already produced 4,600 payers and $820,000 cash in 30 days. These are baseline, not investment value. The full 100,000-person intake is context. Its 92,000 without recent portal visits are not proven never-visitors, reachable SMS users or responders. Expanding to them requires separate discovery.

Eight incomplete-plan interviews and 50 selected support conversations suggest date mismatch, interruption and account questions. They do not establish prevalence. Affordability or creditor terms may dominate. Portal-only guidance is the strongest simpler alternative; cost it alongside conversation and scheduled replacements.

Model progression and incremental cash

The interactive calculator compares current and adaptive rates for opening the portal, starting a plan, confirming it and making a first plan payment. Consumers enter at their existing stage. The starting graduation rates are an illustrative calibration of the earlier +3-point paid-conversion case, not measured transitions or expected AI performance.

At 60% encounter, the illustration gives 1,286 more portal visitors, 1,029 more starts, 463 more confirmations and 324 more first payments. Fees require settled cash. Test substitution and smaller payments: the default assumes all extra payers are incremental at the historical $170/$180/$190 proxy.

Reached paid-lift sensitivity	Extra cash per intake	Annual contribution	Simple payback
+1 point	$18,960	$14,026	42.8 months
+3 points, reference	$56,880	$73,181	8.2 months
+5 points	$94,800	$132,336	4.5 months
+10 points, speculative	$189,600	$280,224	2.1 months

These scenarios assume 15% fees less 2% payment costs, two extra support conversations per 100 reached consumers at $6, net-zero replacement SMS and no live AI cost. Conversation costs remain editable. Annualisation assumes twelve new equivalent intakes. This is a range, not a forecast or confidence interval. Do not add the paid-lift and graduation models together.

The reference leaves only $23,181 after the $50,000 build, before unpriced costs. No improvement loses $15,552 annually before build. Halving marginal cash or genuine incrementality gives 20.8-month payback. The original economics remain a testable, thin investment case. Model.

Scope and AI autonomy

Start with a small approved cohort and intent set, existing approved options, template replies and human assistance. Secure tasks require explicit confirmation. Exclude new creditor terms, autonomous agreements and autonomous payments. A reply, plan start or confirmation earns no collections fee until cash settles.

The demonstration is simulated. Final integration verification owns claims about working capability. No live model, provider or money movement is assumed. AI can help classify stated needs or draft within an approved workflow; it cannot determine legal scope, grant permission, infer affordability, change terms, authorise funds or declare settlement. The current classifier is deterministic. Humans approve US content and production behaviour.

Four-week launch and experiment

Week 1: confirm inbound feasibility, channel scope, identity/security, approved content, support and cost. Name compliance and stop owners. Choose the limited test, simpler alternative or deferral. Week 2: prepare approved conversation and safe routing. Week 3: connect tasks, continuity and measurement. Week 4: test, obtain human compliance approval and launch a capped canary if prerequisites pass.

Measure graduation at the actual entry point. For click-only consumers, focus first on portal arrival. For existing visitors, focus on plan start. For incomplete plans, focus on confirmation. Follow all paths through settled payment, without treating more starts or plans as realised financial value. Returning later, seeking help and declining unsuitable terms remain valid outcomes.

Randomise consumers at t0, stratified by entry cohort, against the existing journey. Primary outcome is treatment-control 30-day net settled cash per assigned consumer, valued at 13% less measured incremental costs. Include non-responders and later suppressed accounts in their original arm. Responders select themselves, so their results alone cannot establish causal lift. Report uncertainty, reach, actual transition rates, support and consumer guardrails. Follow 60/90-day cash and adherence to distinguish acceleration or substitution from durable recovery.

Cash variance and graduation rates are unknown. Size the primary outcome before claiming power. The earlier conversion-power calculation does not power the cash outcome. Complete 30-day windows plus next-day reporting after launch.

Safety and investment gates

Before proactive handoff, re-read current scope, permissions, payment, plan, balance, holds, verified local time and shared contact history. Unknown required state blocks outreach. Preserve disclosures and opt-outs. Apply the supplied weekday window and inclusive one-per-24-hour, three-per-seven-day limits across channels. Prefer replacing a scheduled contact. Silence is not permission for repeated chasing. Requested service replies still require approved classification and safeguards. Humans or authoritative services apply holds; the classifier routes concerns.

Expand only with safe, mature randomised evidence whose conservative contribution estimate supports payback within 12 months at verified reach and costs. Validate the simpler alternative and durability. The four-month planning goal requires $12,500 monthly contribution; the reference supplies $6,098. Use specific graduation improvements and actual costs to assess that gap. A target is not a predicted effect. Build, ramp and reporting delay recovery.

Change if encounter rates are weak, support is overloaded or affordability and terms dominate. Permit one bounded extension if uncertainty spans the investment hurdle. Pause for prohibited contact, disclosure failure, unauthorised terms/payment or serious harm. Investigate complaint increases above 0.10 points or opt-out increases above 0.50 points versus control; these proposed triggers do not establish safety below them. Stop for recurring failures, uncontrolled harm or a precise upper contribution estimate below the hurdle. AI assisted analysis, drafting, code and checks; the candidate owns the recommendation.